



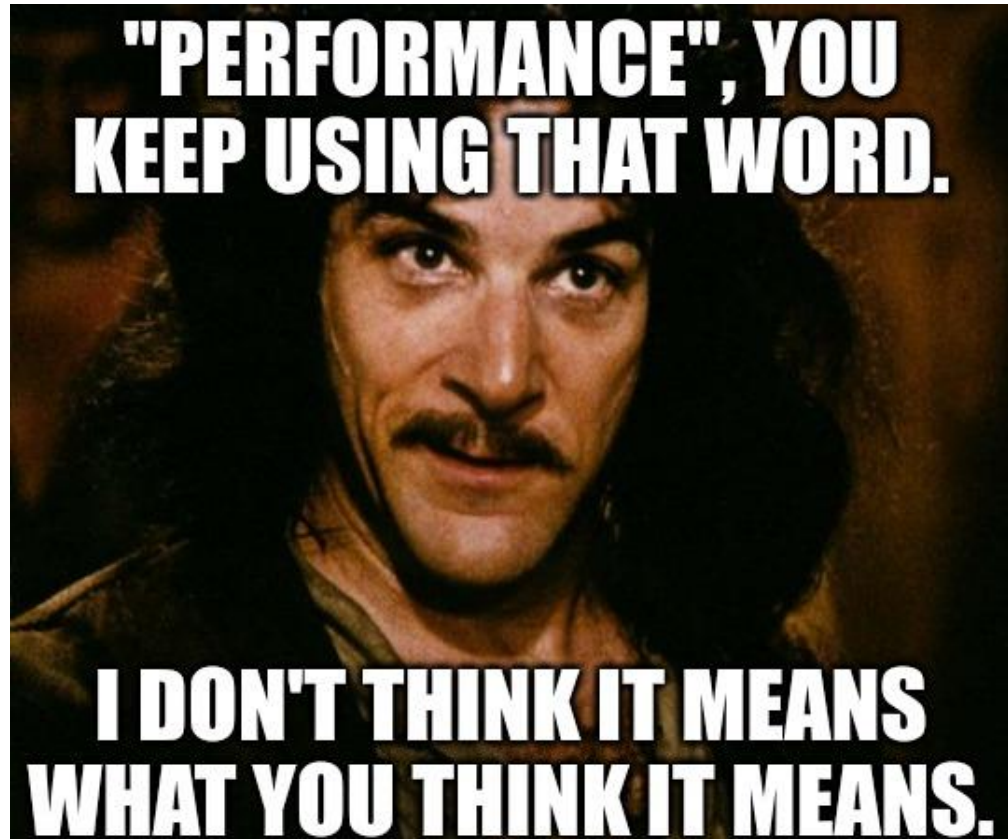


Faster Pandas

miki @tebeka

CEO, CTO, UFO ...





Performance = CPU/Memory

Why?

Why Not?

In [1]:

“Talk is cheap. Show me the code.”
- Linus Torvalds

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel Optane memory access	<10 μ s	7 hrs
NVMe SSD I/O	~25 μ s	17 hrs
SSD I/O	50–150 μ s	1.5–4 days
Rotational disk I/O	1–10 ms	1–9 months
Internet call: SF to NYC	65 ms	5 years
Internet call: SF to Hong Kong	141 ms	11 years

Computer Latency at Human Scale

Culture >> Process

Thank You!

miki@353solutions.com



COURSE

FEATURED

NEW

Faster pandas

By: Miki Tebeka • Released 6 days ago



1h 23m 30s left



COURSE

Faster Python Code



LinkedIn • By: Miki Tebeka • Mar 2018



1h 21m 18s left



Pandas Brain Teasers

Exercise Your Mind



Miki Tebeka
edited by Margaret Eldridge